

data-provisioning

iexec

oracles

scalability

redchainlab

security

cloud-computing

standardization

decentralization

blockchain-trilemma

events

collusion-resistance

interoperability

Matthieu BETTINGER

About me

Matthieu BETTINGER

- INSA-IF 2021
- Double-degree with Uni. Passau

2021 projects:

- “*Collusion-Resistant Worker Set Selection for Transparent and Verifiable Blockchain-Based Voting*” with Lucas Barbero & Omar Hasan
- “*Blockchain interoperability – Cross-chain oracles with Decentralized Cloud Computing*” with iExec, INSA Lyon & Uni.Passau

2022–2024:

- PhD thesis on “*Blockchain interoperability & composability*” in RedChainLab

Summary

- **iExec Decentralized Cloud Computing**
- **Using blockchains to secure applicative protocols**
 - “*Collusion-resistant worker set selection*”
- **RedChainLab**
- **Oracles: Connecting blockchains & the outside world**
 - Theory & Practice
 - Standardization & Interoperability
 - Data provisioning trade-offs
 - “*Oracles with Decentralized Cloud Computing*”
- **Project/Practical work**

Goals

- Real-world applications of blockchain/smart contract technology, how they secure protocols
- Trustworthy data provisioning issues and countermeasures
- Event-based programming with Solidity/Web3J(S)

A Trusted, Decentralized Marketplace for Cloud Computing

By **Anthony Simonet-Boulogne, PhD**

Research Scientist, iExec Blockchain Tech

Revisited by **Matthieu Bettinger**



INSA Lyon
Dec. 7th 2021

What we do



iExec provides an **open** and **decentralized** cloud computing **marketplace**.

Connects **providers** with **users**: anyone can trade **computing power, datasets, and applications**.

No need to trust iExec or anyone else: just **trust the blockchain** and the code.

Blockchain & Decentralisation

A decentralised, immutable and verifiable digital ledger consisting of transaction records distributed across many computers.



Decentralised



Immutable



Verifiable

Smart contracts?

Self-executing and self-enforcing programs that can read and write the state of a blockchain.

- ✓ Transparent
- ✓ Auditable
- ✓ Autonomous
- ✗ Hard to program
- ✗ Limited in size
- ✗ Often extremely critical
- ✗ **Cannot access external data**

```
contract Coin {  
    address public minter;  
    mapping (address => uint) public balances;  
  
    function Coin() public {  
        minter = msg.sender;  
    }  
  
    function mint(address receiver, uint amount) public {  
        if (msg.sender != minter) return;  
        balances[receiver] += amount;  
    }  
  
    function send(address receiver, uint amount) public {  
        if (balances[msg.sender] < amount) return;  
        balances[msg.sender] -= amount;  
        balances[receiver] += amount;  
    }  
}
```

Blockchains

What they are

- A way for arbitrary people and organisations to collaborate without having to trust anyone.
- A tool for transparency and democracy.
- A platform for deploying unstoppable programs.

What they (usually) are not

- ✗ Fast
- ✗ Cheap
- ✗ Easy to program
- ✗ User friendly

But we're working on it ;-)

iExec history

Founded in 2016 by Gilles Fedak (Inria) & Haiwu He (Chinese Academy of Sciences)

April 2017: ICO raised 10,000 Bitcoins within 3 hours

Based in Lyon



A not so new idea...

*We described a computational model based upon the classic science-fiction film, **The Blob: a program that started out running in one machine, but as its appetite for computing cycles grew, it could reach out, find unused machines, and grow to encompass those resources.** In the middle of the night, such a program could mobilize hundreds of machines in one building; in the morning, as users reclaimed their machines, the “blob” would have to retreat in an orderly manner, gathering up the intermediate results of its computation. (This affinity for night-time exploration led one researcher to describe these as “vampire programs.”)*

(John F. Shoch and Jon A. Hupp, 1982)

iExec allows individuals and enterprises to monetize their computing power, applications and datasets.

**Computing
power**



Applications



Datasets



Why decentralise the cloud?

Centralized Computing

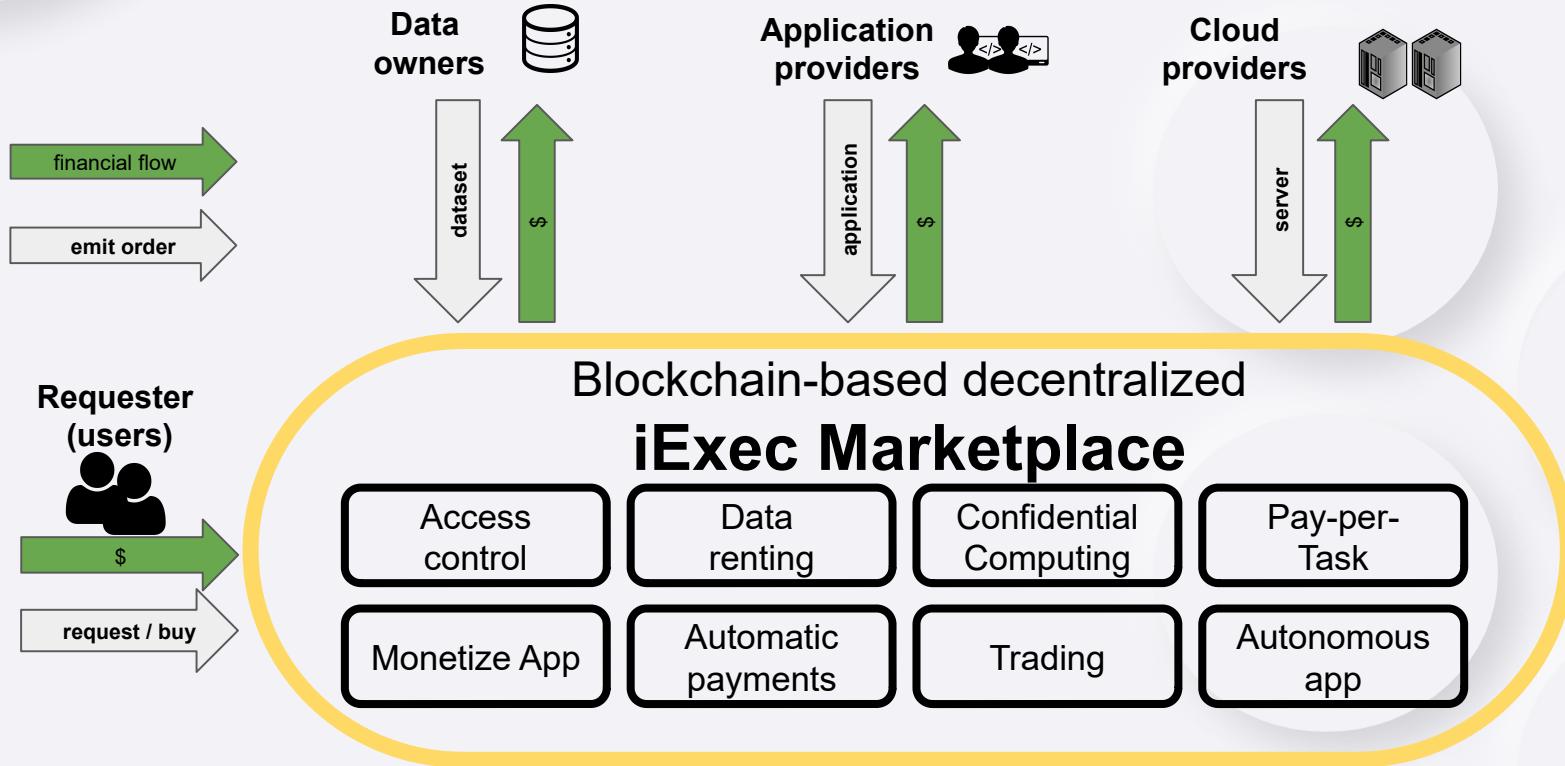
- ✗ Unfair pricing
- ✗ Vendor lock-in
- ✗ Limited transparency
- ✗ Limited accountability
- ✗ No provenance information
- ✗ Possible censorship

Decentralised Cloud Computing

- ✓ Market-based prices
- ✓ Fair competition between providers
- ✓ Smooth business agreements
- ✓ Complete execution history on the blockchain
- ✓ *Unstoppable* marketplace: censorship is impossible



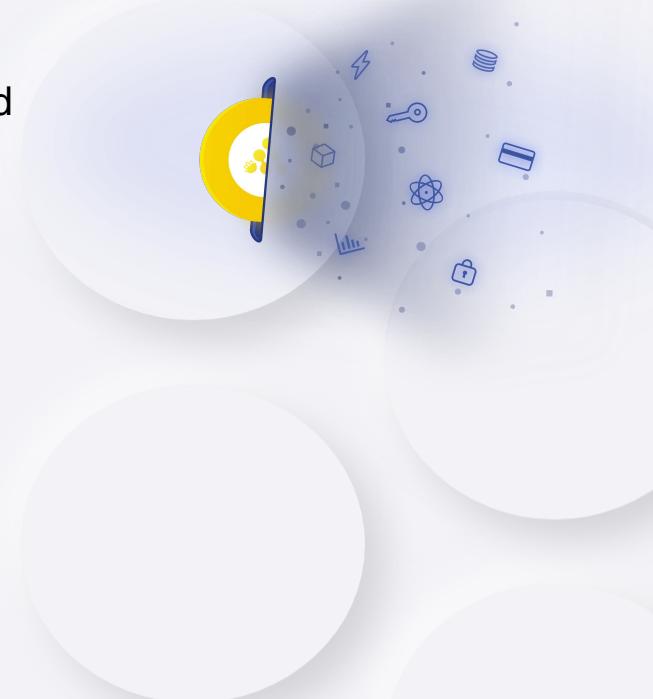
iExec overview



The iExec token: RLC

RLC is an ERC-20 utility token.

- RLC is necessary to access the iExec decentralised cloud
- Providers are paid with RLC
- RLC allows to build incentives in the network
- RLC creates a specific market for cloud computing



Two types of tasks

with configurable confidence and privacy

Standard tasks

Run on untrusted resources, delegate trust to the blockchain

- Replication level depending on desired confidence
- Decentralized consensus
- On-chain reputation
- Staking & economic incentives
- Deterministic

TEE tasks

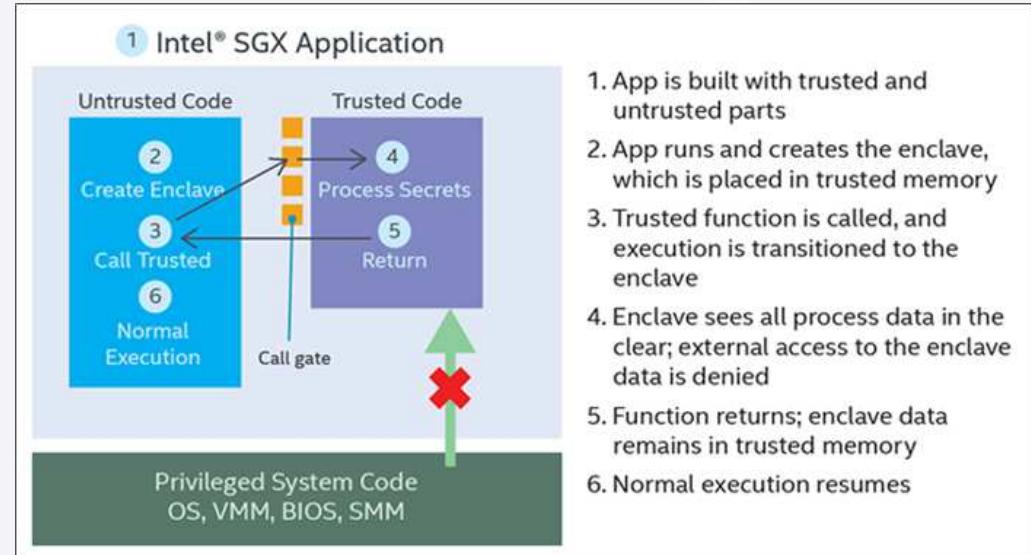
Run isolated within an Intel SGX TEE (Trusted Execution Environments)



- End-to-end encryption of data & result
- Enclave attestation proves that the task was run in TEE
- Result signature with enclave key: no need for replication
- Determinism not required

Trusted Execution Environments?

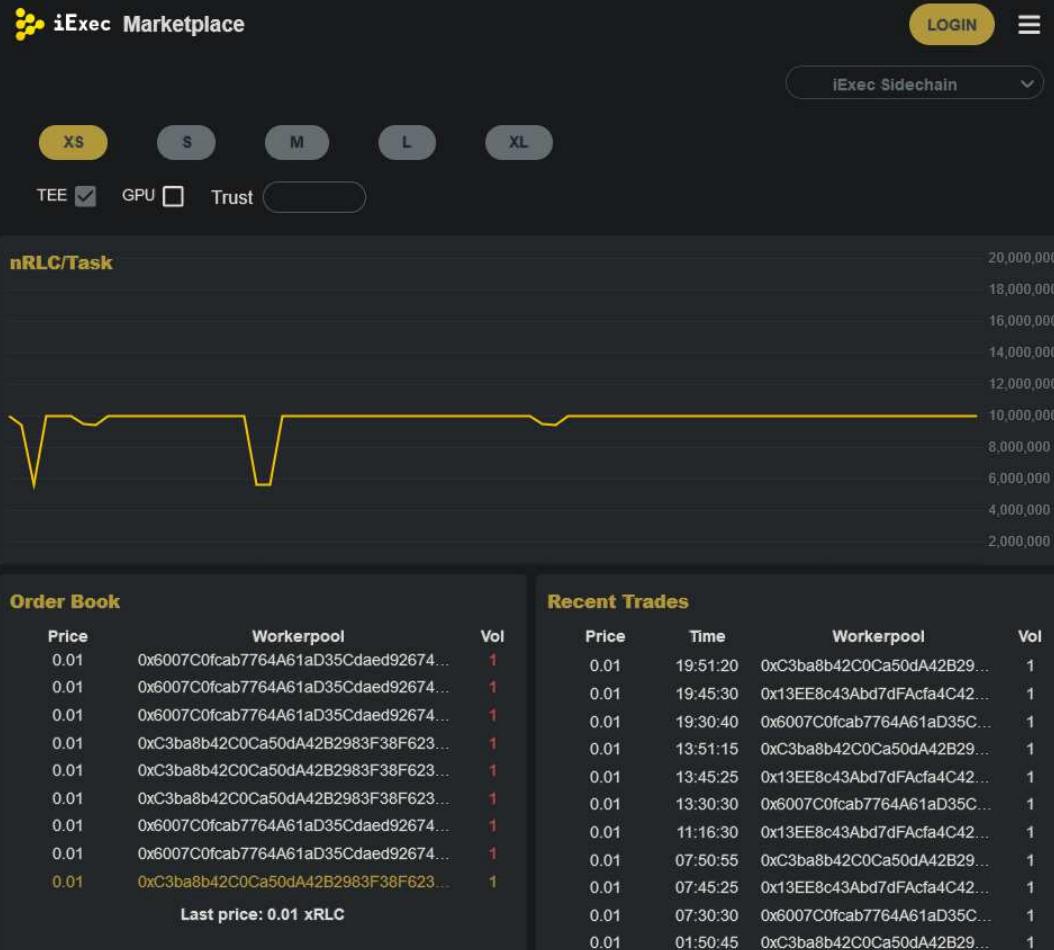
- Secure part of a CPU with encrypted memory space
- Memory & Code protected from host (even root)
- Hardware based security (private key in silico)
- Can be remotely attested
- (Soon) available on hardware from various vendors



Intel® Software Guard Extensions application execution flow.

Components: iExec hub

- Managed by Smart contracts
- Repository of registered resources (dApps, workerpools, ...)
- Storage of task results and metadata
 - Task specifications
 - Execution details
 - Off-chain storage link



Components: Workerpool

- Composed of a scheduler and multiple workers
- Incentives:
 - Staking
 - Reputation
- Scheduler objective:
 - Listen to incoming work
 - Distribute work fairly among workers
- Worker objective:
 - Correctly execute tasks given by the scheduler

Components:

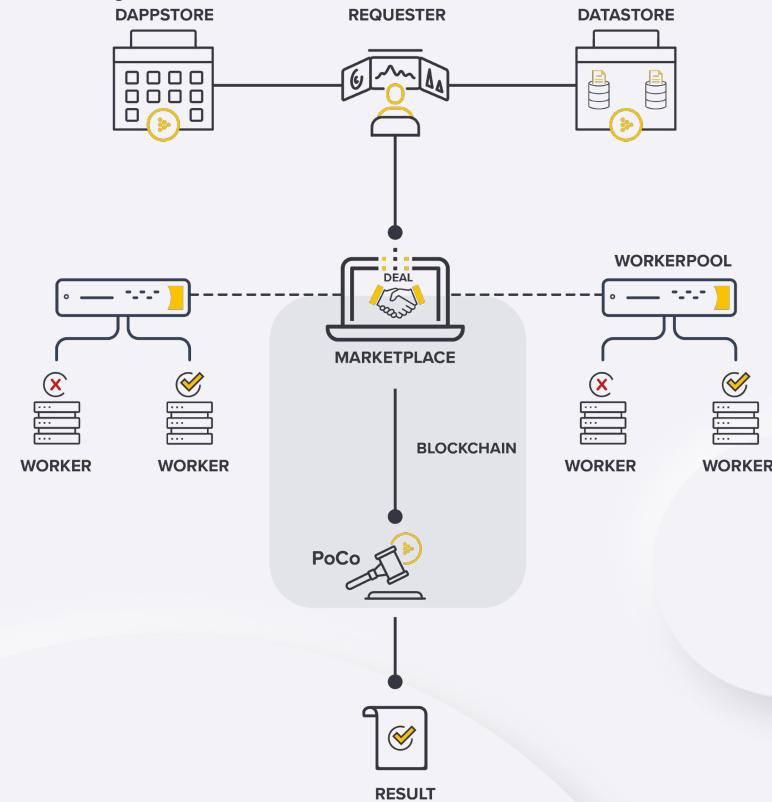
Secret Management Service

- Secured inside in a Trusted Execution Environment (not yet, but in V7)
- Stores secrets of stakeholders:
 - Dataset decryption keys
 - Input files decryption keys
 - Output files/Results encryption keys
- Attests TEE workers

Proof-of-Contribution (PoCo)

On-chain validation than an off-chain task was performed correctly

1. One task = 4 orders, signed off-chain with an Ethereum wallet:
 - **apporder** signed by the developer
 - (**datasetorder** signed by the dataset provider)
 - (**workerpoolorder** signed by a worker pool scheduler)
 - (**requestorder** signed by a requester)
2. Orders are matched on-chain: [poco.matchOrders\(\)](#)
(Check signatures, parameters, balances, ...)
3. PoCo seals a deal & workers start computing
4. Workers send result hash back to PoCo
5. PoCo compares results, manages reputation, triggers payments.



Adjusting trust: Sarmenta Voting

- Deterministic execution replicated **N** times: only one correct result exists
- Given **R** distinct results ($1 \leq R \leq N$):
 - For each result **r**, obtained by **n** among **N** workers:
 - Probability that **r** is correct and all others are false
- User-defined threshold **t**:
 - If no proposed result has $Crt(r) > t$
 - Then more workers are dispatched

$$Cr_t(r) = \frac{\tilde{P}_t(r)}{1 + \sum_{h \in R_t} \tilde{P}_t(h)}$$

https://github.com/iExecBlockchainComputing/PoCo/blob/v5/audit/docs/iExec_PoCo_and_trustmanagement_v1.pdf

iExec Research Projects

H2020 ONTOCHAIN

Building an ecosystem for trustworthy content handling & information exchange



Keywords: Semantic Web, Oracles, Decentralized Identities, integration, applications

2020–2023

H2020 DATA CLOUD

Enabling the Big Data Pipeline Lifecycle on the Computing Continuum



Keywords: Fog/Edge Computing, Big Data pipelines, self-* cloud computing, Industry 4.0

2021–2024

ANR RedChainLab

Scalable, trusted and privacy preserving decentralized marketplaces

Joint laboratory between the **DRIM research team** (LIRIS, CNRS) and **iExec**

Keywords: blockchain, decentralized cloud computing, edge computing, security, TEE, Federated Learning

2021–2024

Try us!

<https://developers.iex.ec/>

<https://iex.ec/grants/>



Join us!

<https://iexec.flatchr.io/>

[https://gitcoin.co/explorer?
network=mainnet&idx_status=open&applicants=ALL&key](https://gitcoin.co/explorer?network=mainnet&idx_status=open&applicants=ALL&key)

Anthony Simonet-Boulogne
anthony@iex.ec

Collusion-Resistant Worker Set Selection for Transparent and Verifiable Blockchain-Based Voting

Matthieu BETTINGER
Lucas BARBERO
Omar Hasan

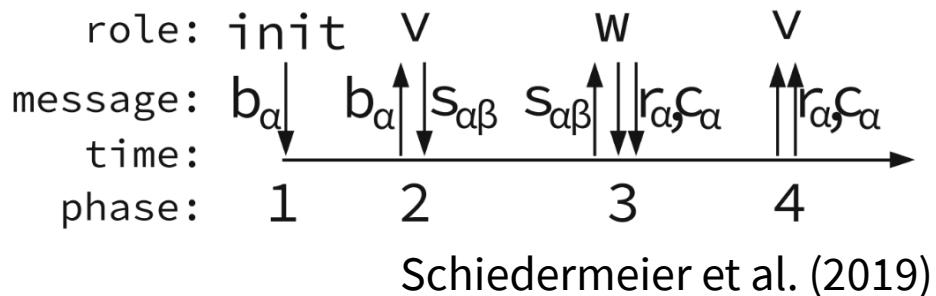
Context

- Schiedermeier's blockchain-based voting protocol (2020)
 - P participants, and W workers chosen among them
 - Workers compute the referendum's result
 - Using SMPC:
 - Shamir's Secret Sharing scheme
 - Homomorphic encryption
 - All messages are recorded on-chain

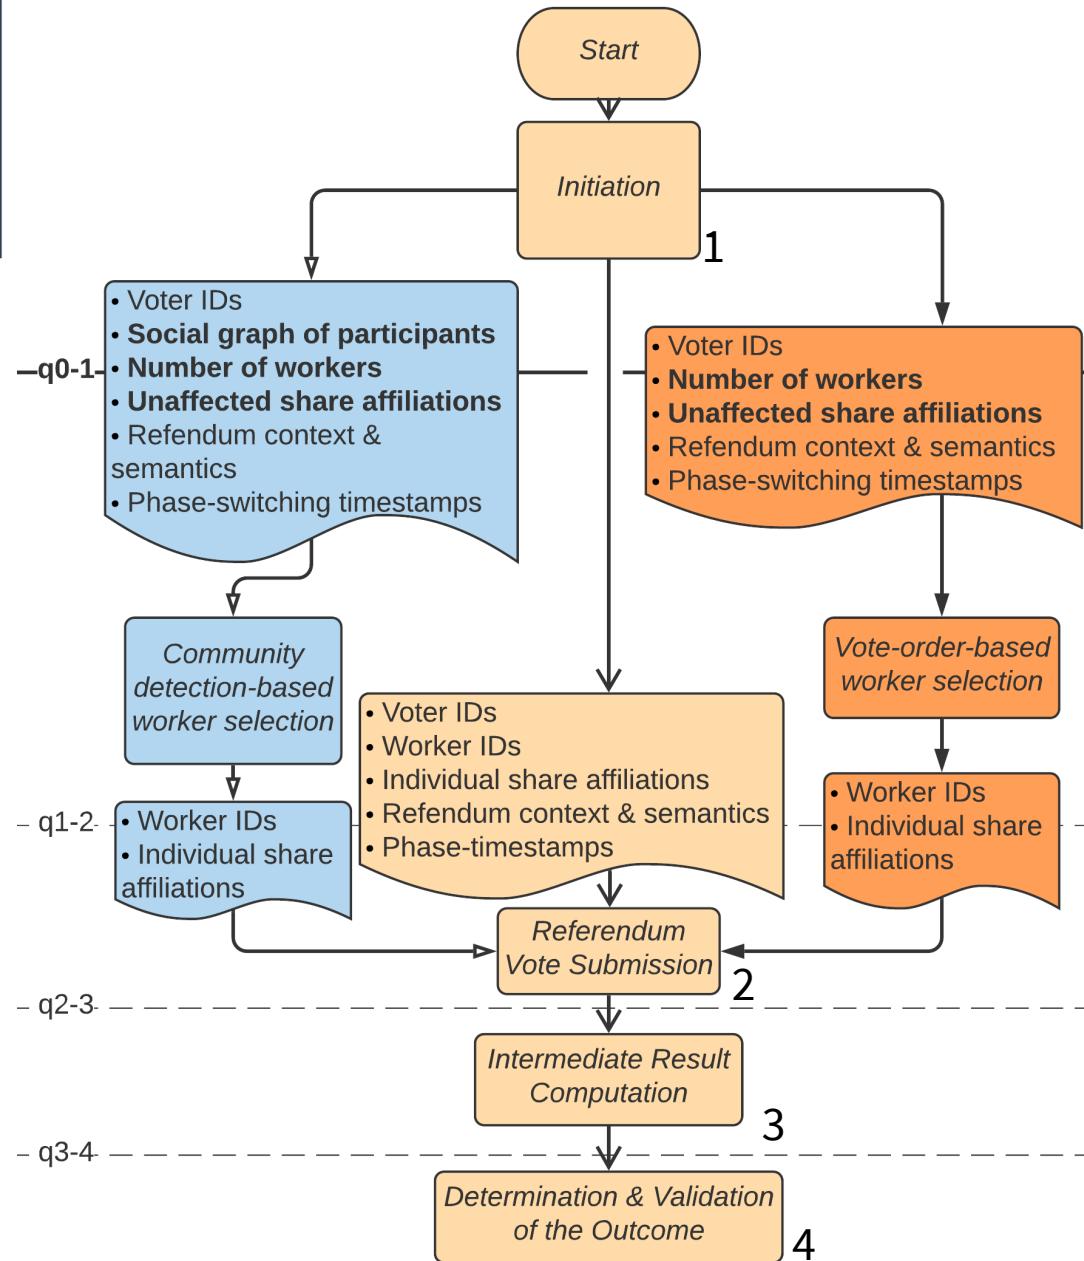
Context

- However: workers are chosen by a single trusted entity
→ Risk of collusion
- Goal:
→ Collusion-resistant & verifiable selection of W workers

Process



Process integration diagram of proposed solutions



Bettinger et al. (2021)

Proposal 1:

Verifiable random number generation

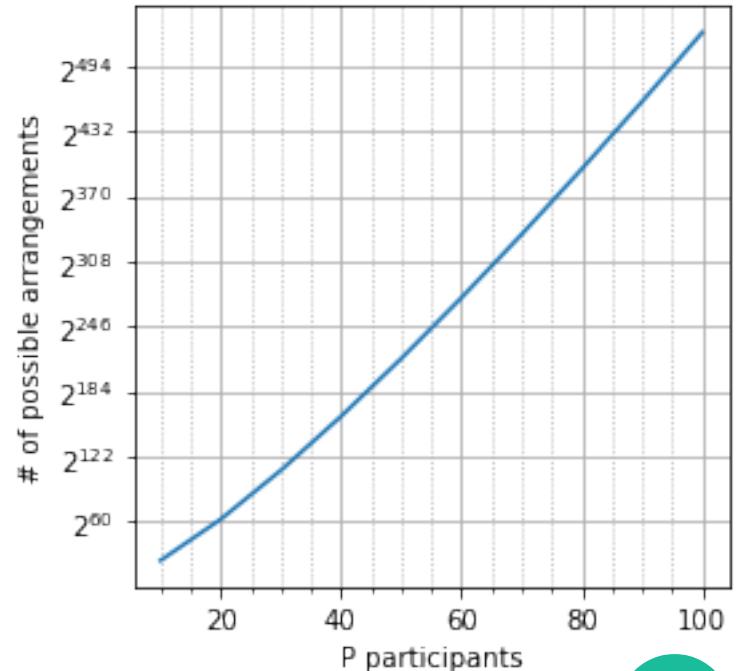
Table 1 Three methods to attribute a number to a list V of comparable non-reoccurring elements (integers, character strings,...) knowing its superlist P .

Numbering Method	Expression	Output Space
Permutation $PN : V \rightarrow \mathbb{N}$	$\sum_{i=0}^{ V -1} (i! \sum_{k=0}^{i-1} \mathbb{1}_{x_k < x_i})$	$\llbracket 0; V ! \rrbracket$
Combination $CN : V, P \rightarrow \mathbb{N}$	$\sum_{k=0}^{j(0)-1} \binom{ P -j(0)+k}{ V -1} + \sum_{i=1}^{ V -1} \sum_{k=0}^{j(i)-j(i-1)-2} \binom{ P -j(i)+k}{ V -i-1}$	$\llbracket 0; \binom{ P }{ V } \rrbracket$
Arrangement $AN : V, P \rightarrow \mathbb{N}$	$\sum_{i=0}^{ V -1} \left(\frac{ P !}{(P -i)!} + CN(V, P) * V ! + PN(V) \right)$	$\llbracket 0; \sum_{v=0}^{ P } A_v^v \rrbracket$

- **PN : Ordering of V voters**
- **CN : Which subset of V participants voted**
- **AN : PN & CN + Sum(AN with fewer voters)**

10 participants : 10^6 possibilities

100 participants : 10^{158} possibilities

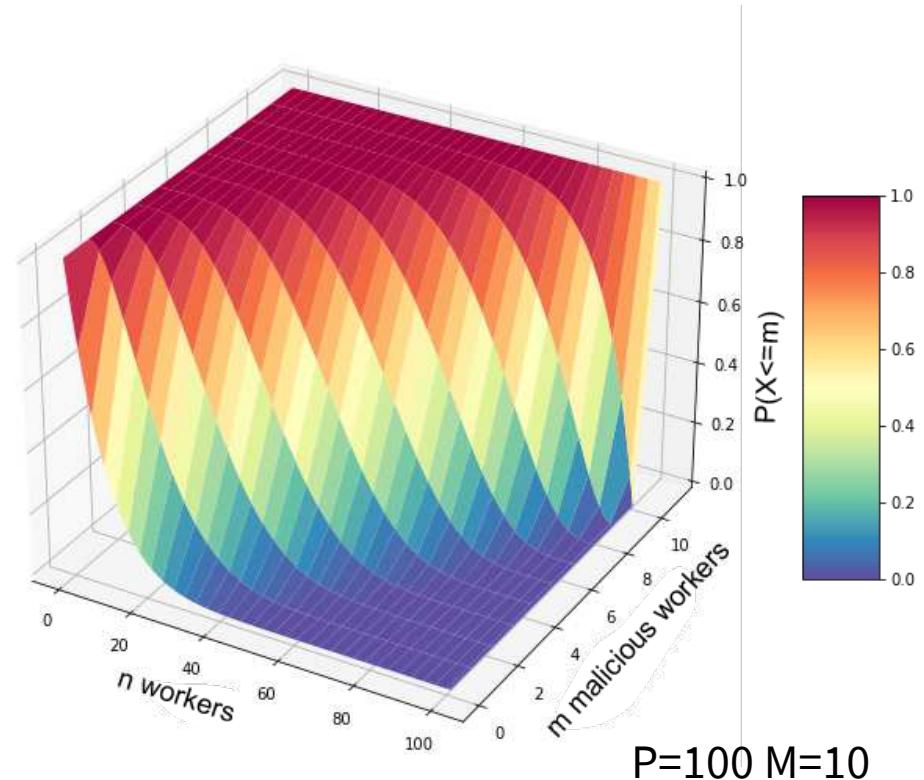


Proposal 1: Probability Distribution

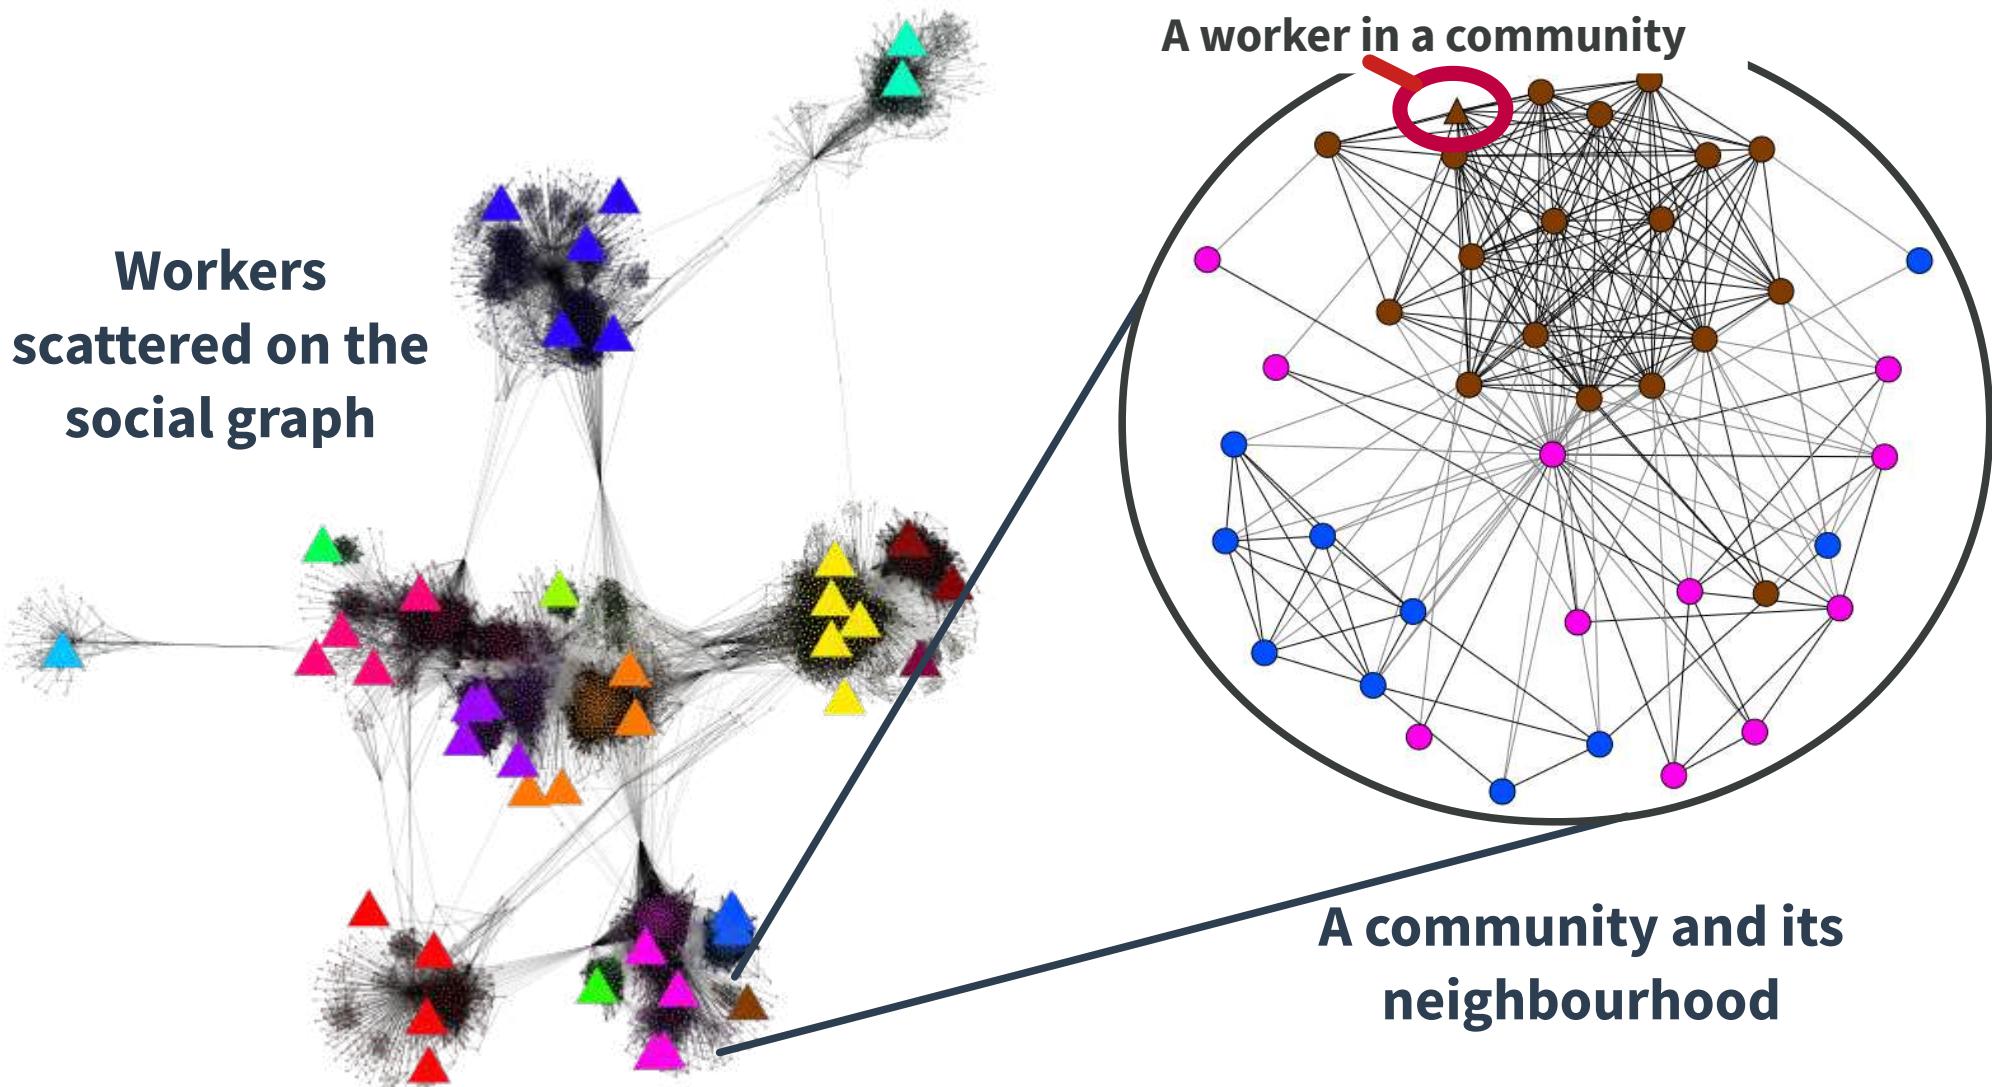
- Hypergeometric Law

$$H(n, M/P, P)$$

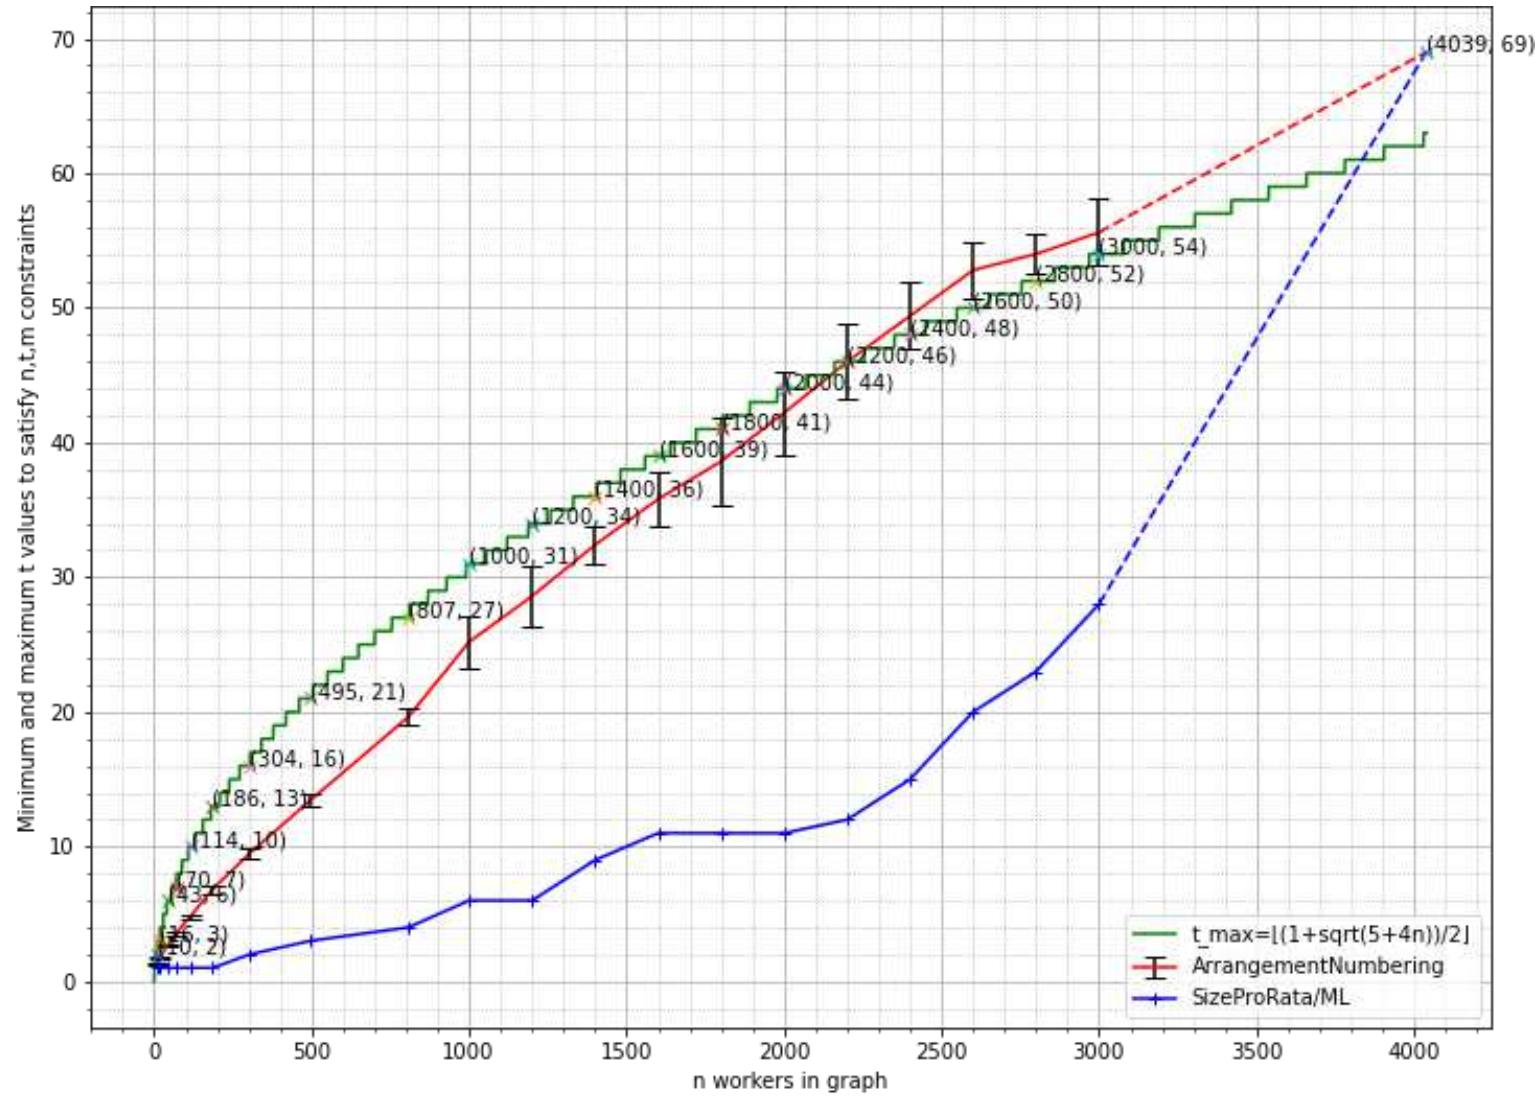
- P participants
- $M \leq P$ malicious participants
- n workers
- m malicious workers



Proposal 2: Worker distancing



Results: Cliques of adjacent workers



Application to iExec's worker selection

- Scheduler: trusted to distribute work fairly
 - Proposed solutions can be used
- Difference with Schiedermeier's:
 - iExec workers perform multiple tasks
 - versus anonymized participants for each vote

Application to iExec's worker selection

- Aral et al. (2020)
 - Use of workers' task execution history
 - Clustering of workers that fail together
- **Algorithm to maximize success probability**

RedChainLab

iExec – DRIM

Matthieu BETTINGER

RedChainLab

- **Joint laboratory**
 - Between the LIRIS-DRIM team and iExec
 - Coordinated by Sonia BEN MOKHTAR (BFT, Distributed Systems)
- **From 2022 to 2025**

~2 end-of-studies projects per year

Research fields

- Blockchain Scalability
- Trusted Execution Environments
- Confidentiality
- Genericity/Interoperability/Composability

Oracles between blockchains and the outside world

Matthieu BETTINGER

Summary

- Motivation
- Oracles: Theory and practice
- Standardization & Interoperability
- Data provisioning trade-offs

Motivation

- Blockchains are closed systems
 - they cannot read/write data of off-chain systems
 - off-chain entities provide data
 - Need for data provisioning systems
- Use cases for smart contracts in IoT (Industry 4.0, Smart Cities,...), Education, Health and more
 - Relying on critical off-chain data or events
 - Need for **trustworthy** data provisioning systems

Oracles

- Link blockchains to off-chain data and events (even to other blockchains) and vice-versa
- Two components:
 - Off-chain: data collection, processing, transmission wrt. requests
 - On-chain: interfaces requesting/accepting data + storage
- The proportion of on/off-chain computation depends on the architecture

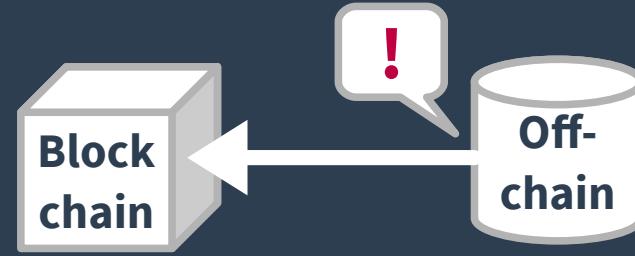
Types of oracles (Mühlberger et al. 2020)

- Is it the data recipient who requests the data transfer ?
 - Yes → Pull-based
 - No → Push-based (another entity is responsible)
- What is the data transfer path ?
 - Blockchain to outside world → Outbound
 - Outside world to blockchain → Inbound

Example for showcasing oracle types

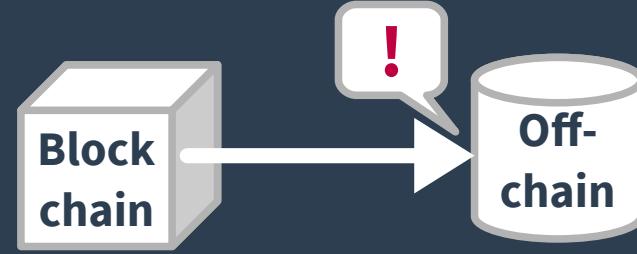
- **Blockchain-based betting platform**
 - Users populate events (e.g. sports matches) on-chain
 - Any interested user can lookup current betting opportunities and bet on them
 - The platform requests the results from a trustworthy oracle system
 - Once a result is known, the platform broadcasts winners of the corresponding bet

Push-inbound oracles



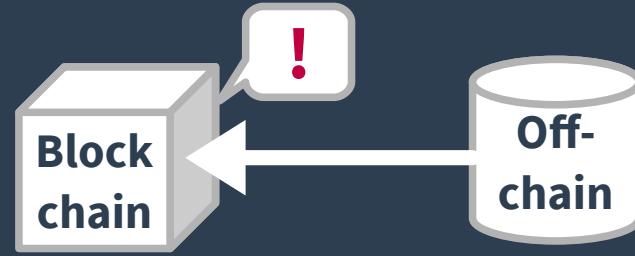
- Goal: Feed relevant off-chain data to the smart contract
- Behaviour:
 - Off-chain states/events are monitored
 - EOAs (Externally Owned Accounts) send transactions
 - Passive recipient smart contract
- Example:
 - Add a sports event on which to bet in the platform

Pull-outbound oracles



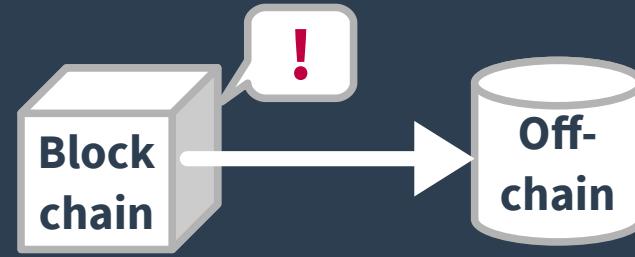
- Goal: Retrieve relevant data from the smart contract
- Behaviour:
 - Queries: off-chain systems read the state of the contract
 - Passive data source smart contract
- Example:
 - Lookup current standings for a bet

Pull-inbound oracles



- Goal: Receive relevant off-chain data
- Behaviour:
 - Events (+ requests stored in the smart contract's state)
 - Active smart contract (still needs an EOA to initiate the request's creation)
- Example:
 - Request the result of a sports match (if available)

Push-outbound oracles



- Goal: Signal relevant chain data to offchain systems
- Behaviour:
 - Events (+ requests stored in the smart contract's state)
 - Active smart contract (still needs an EOA to initiate the request's creation)
- Example:
 - Broadcast winners of a bet

Real-world oracle systems

- **ChainLink V1: On-chain data aggregation**
 - Provides predetermined feeds of temporal data (e.g. ETH-USD exchange rate)
 - Off-chain nodes periodically fetch new values from an off-chain source
 - Each node pushes its findings on-chain
 - Values are aggregated on a given metric (avg, med, ...)
 - Outlier values' nodes are punished
- **Pull-inbound oracle (for periodic updates)**
- **Push-inbound oracle (for value-adaptative updates, e.g. based on relative change)**



Event-based programming

- In Solidity:

`event NamedEvent(Type1,T2,T3,...);` → declaration in class' body

`emit NamedEvent(param1,p2,p3,...);` → call in a function's body

- Event parameters (up to three) can be *indexed* for easier filtering

- In Java/web3j:

`contractInstance.NamedEventFlowable(...).subscribe(event → { /*event handling*/ })`

- In JS/web3js:

```
var namedEvent = constructInstance.NamedEvent();
```

```
namedEvent.watch(function(error, result){ /*event handling*/ });
```

Standardization

- EIPs: Ethereum Improvement Proposals
 - EIP-1154: push-inbound & pull-outbound oracles
 - Abandoned by the community (but still used by iExec)
 - EIP-2362: interface for pull-based price-feed oracles
 - By the Alliance of Decentralized Oracles
 - In reaction to the unpopularity of EIP-1154

Interoperability

- Increasing difficulty
 - Standard < Manual interop. < Automated interop.
- Interoperating entities:
 - smart contracts (SCs) in the same blockchain
 - SCs and oracle systems
 - SCs across distinct blockchains (through oracles)

On-chain VS Off-chain computations

- Oracle systems have both on- & off-chain parts
- On-chain:
 - Greater security/decentralization → Greater costs/delays
- Off-chain:
 - Less costs but loss of blockchain properties

→ Balance of costs and security to be found

On-chain VS Off-chain computations

- **ChainLink V2: Off-Chain Reporting**
 - Rounds of an off-chain BFT algorithm
 - A selected leader generates the new value to report
 - Other nodes agree by signing the report
 - The agreed-upon report is pushed on-chain



OFF
+
ON

I
ON
+
OFF
T
ON
I

Other data acceptance mechanisms

- **Astrea: Voting**
 - A proposition is submitted on-chain
 - Participants check whether the proposition is true or false
 - Participants vote whether the proposition is true or false
 - After a timeout or sufficient votes, voting is closed
 - The minority loses its staked assets



Other data acceptance mechanisms

- DOS (Decentralized Oracle System): Random selection
 - N participants are randomly dispatched in small groups of size n
 - All nodes in a group compute the result of the same task
 - t or more out of n must sign the same result to attest it
 - This signed result is pushed on-chain
 - Only signers are rewarded

Blockchain interoperability

**Cross-chain oracles
with
Decentralized Computing**

Matthieu BETTINGER

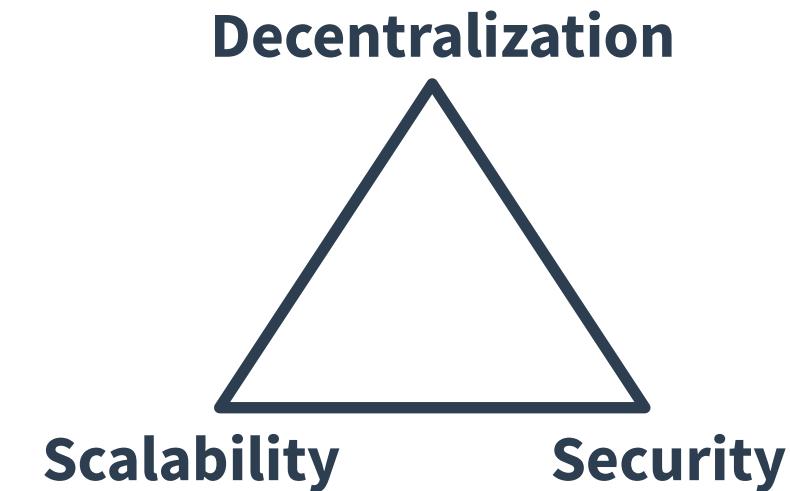
Summary

- Context & Motivation
- State of the Art
- Proposed solutions
- Results & Evaluation
- Conclusion & Future work

Context & Motivation

Blockchain trilemma:

- No "silver bullet"
- Complex user needs
 - Integration of distinct blockchain solutions
 - Oracles linking blockchains
- User needs apply on oracle specs as well
 - Need flexible oracle architectures



Context & Motivation

- Generic oracle use-cases
 - Data queries (request-handling)
 - Triggers (event-handling)
 - Cross-chain calls
 - Non-functional properties
 - Data integrity & availability
 - On/Off-chain verifiability
 - Scalability (throughput, costs)
-
- The diagram illustrates the relationship between the use-cases and properties. A bracket groups the first three items under 'Generic oracle use-cases' and points to the label 'Cross-chain oracles'. The fourth item, 'Cross-chain calls', is shown with an arrow pointing to the label 'Cross-chain smart contracts'.

Context & Motivation

- Variety of use-cases and non-functional requirements by end-users
 - For blockchain systems
 - For oracle systems
- Use of decentralized cloud computing infrastructures for flexibility

State of the Art

- **Datafeeds:**
 - Feed data to blockchains
- + Support all oracle use cases
- Rigid architectures

State of the Art

- **Cross-chain bridges**
 - Relay information between chains
- + **Scalability, data integrity & availability**
- Specialized in requested generic data transfers, no user-defined logic
- (usually) Single data acceptance mechanism

State of the Art

- **Decentralized cloud computing infrastructures:**
 - Execute arbitrary provided logic: a decentralized app
- **+ Sandboxed architectural freedom**
- **- Administrative & Genericity-related overheads**
- **Objective:**
 - Enable flexible datafeeds & bridges with generic DCCI

Context & Motivation

- Generic oracle use-cases
 - Data queries (request-handling)
 - Triggers (event-handling)
 - Cross-chain calls
 - Non-functional properties
 - Data integrity & availability
 - On/Off-chain verifiability
 - Scalability (throughput, costs)
-
- The diagram illustrates the relationship between generic oracle use-cases and their applications. A bracket groups 'Data queries (request-handling)', 'Triggers (event-handling)', and 'Cross-chain calls' under the heading 'Cross-chain oracles'. Another bracket groups 'Cross-chain oracles' and 'Cross-chain smart contracts' under the heading '→ Cross-chain smart contracts'.

Research questions

- **Secure data transfers using DCCIs**
 - data integrity, on/off-chain verifiability
- **Flexibility of oracle system properties**
 - data acceptance, scalability, decentralization
- **Cost reduction strategies**

Proposal 1: Proxy tasks

+ Cost reduction:

- Blockchain trust algorithm (PoCo) delegated on a fast and cheap chain
- Storage reduction: Only results and traceability metadata stored on the costly chain

+ Code & Data integrity: TEE to securely generate & transfer results

Inherited from iExec infrastructure:

- Usability; Storage costs: Inefficient untyped blob storage for results
- Reactivity; Throughput: Single result per task

Proposal 2: Work passes

Basics

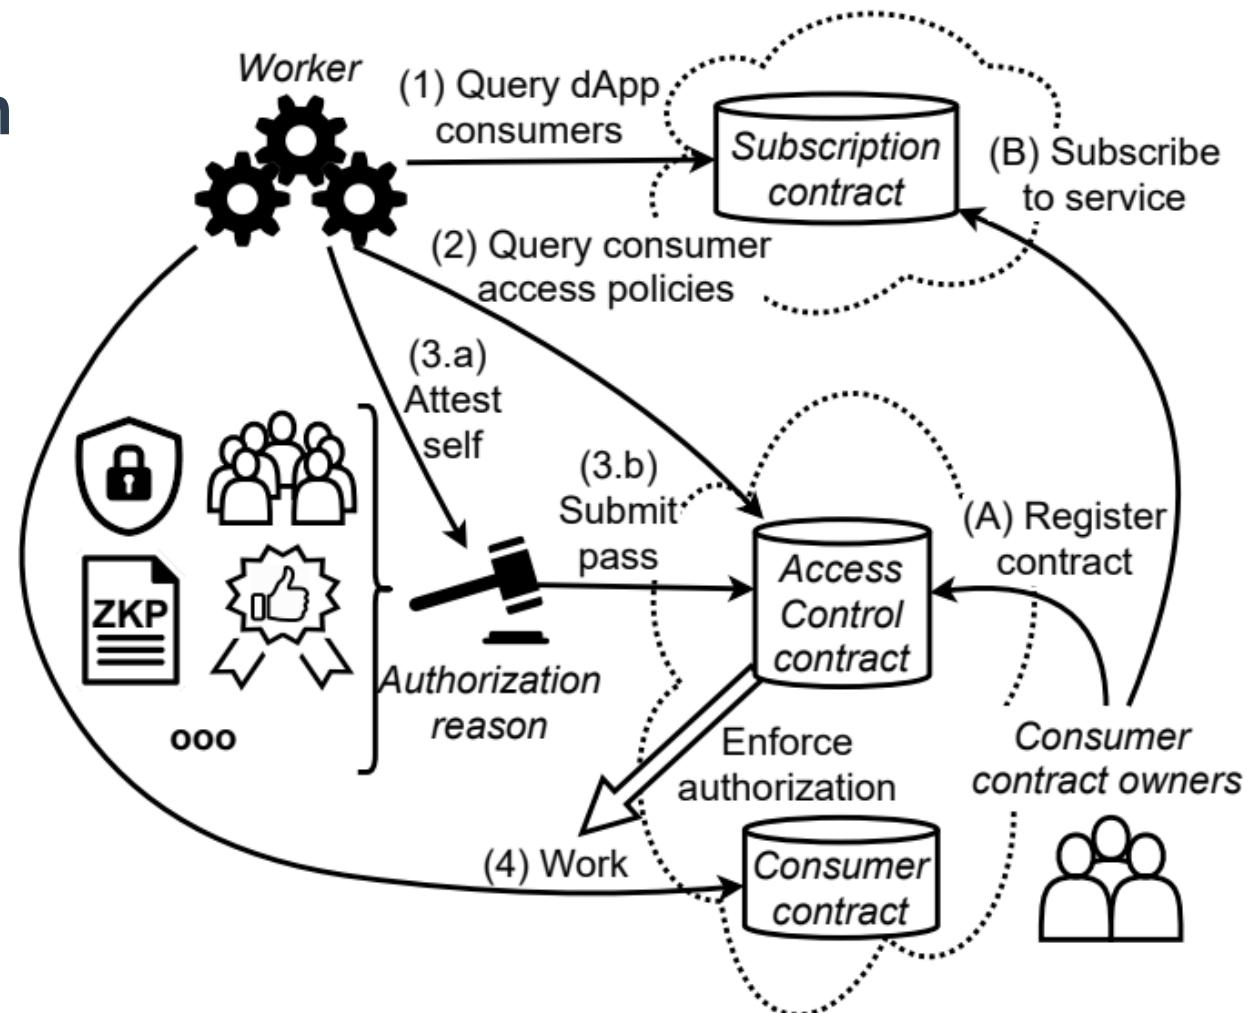
Access control lists (RBAC):

- A pass links:
 - 1 off-chain entity (e.g. an iExec worker)
 - 1 contract (or a subset of methods)
 - 1 scope of authorization
 - OWN > DELEGATE > WORK > NONE
 - For a limited time

Proposal 2: Work passes

Workflow

- Authorization reason
 - Attested TEE
 - Committee consensus
 - (ZK) Proof
 - Trusted identity
 - ...



Proposal 2: Work passes

Protocol properties

- + *Reactivity; Throughput*: Multiple results per dApp execution
- + *Flexible security*: Consumer-defined data acceptance logic
- + *Usability; Storage costs* : Application-specific typed storage
- *Security risk*: Direct access to consumer methods
 - Intermediate dApp-specific storage as buffer
- *Storage costs*: A pass per consumer contract

Proposal 2: Work passes

Protocol optimizations

- One-time passes
 - + No persistent storage costs
 - Repeated authorization verification
- Per-dApp passes
 - 2 passes on-chain : Contract-Service & Service-Worker
 - + Only one pass per worker per dApp per chain
 - Additional storage costs + pass linking costs

Results:

General Metrics

- Full PoCo task execution with TEE worker :
 - ~5 minutes of non-productive processing

Proxy tasks VS current iExec tasks :

- -64 % gas costs

Results :

Algorithmic complexities

Complexity \ Method	Proxy tasks	Time-bounded work passes	One-time work passes	Per-service work passes
On-chain	$O(R * C)$	$O(C * N)$	$O(C * N)$	$O(C)$
Off-chain	$O(R)$	$O(C * N)$	$O(R * C * N)$	$O(1)$

Table 6.2: Administrative cost complexities

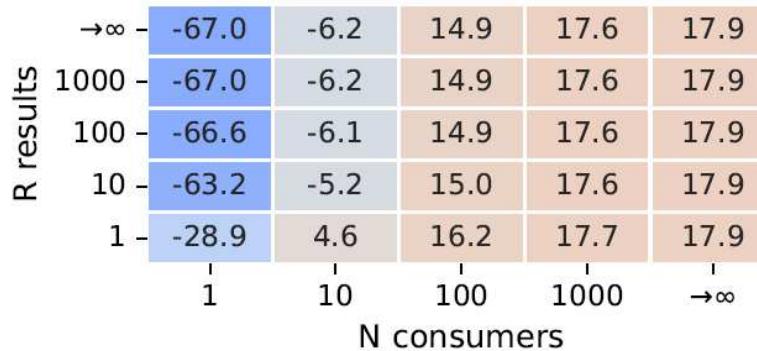
Complexities are relative to the number C of target chains, N consumer contracts per chain (on average), and R results to transfer.

Results:

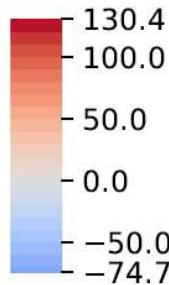
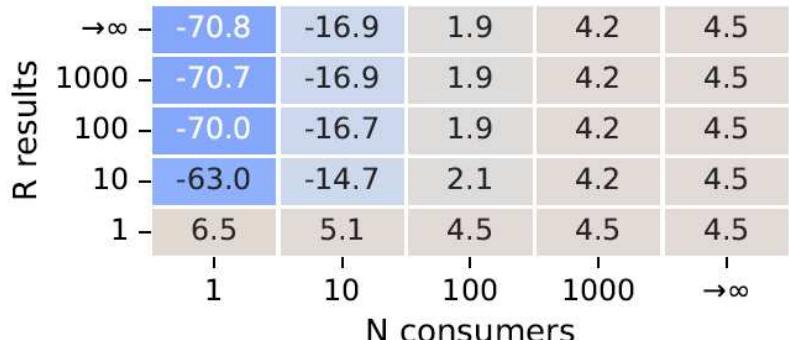
Scaling costs



(a) Time-bounded work passes



(b) One-time work passes



(c) Per-service work passes

$$\begin{aligned}
 & \lim_{N \rightarrow \infty, R \rightarrow \infty} \left(\frac{G(N, R)_X - G(N, R)_{PT}}{G(N, R)_{PT}} \right) \\
 &= \frac{G_{receiveResultX} - G_{receiveResultPT}}{G_{receiveResultPT}}
 \end{aligned}$$

Table 6.3: Relative costs (%) of *work pass* methods compared to *proxy tasks*, depending on the number N of consumers and R results to transfer

Results: Impact of N-Batching

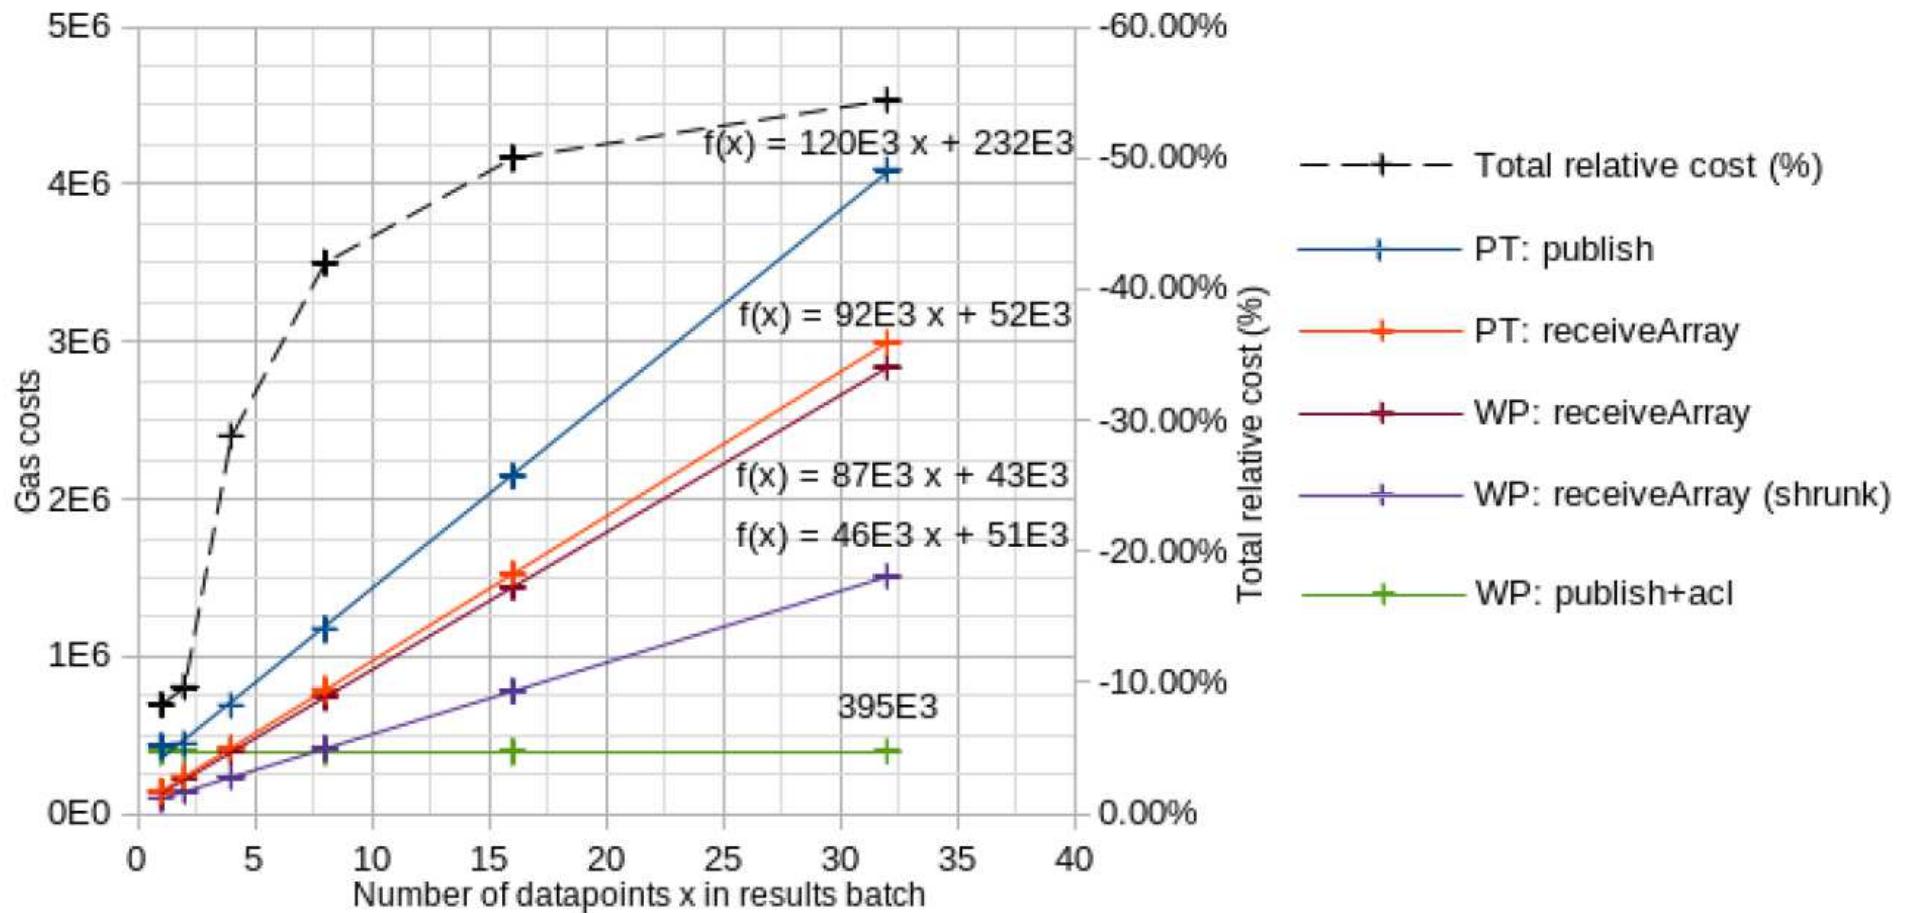


Figure 6.4: Gas costs to transfer n data points at once, using *Proxy tasks* (Proxy) or *Work Passes* (WP)

Conclusion

- Reuse of trusted security mechanisms
- Less storage and processing on costly blockchains
- More genericity:
 - Data acceptance mechanisms
 - Result size and number

Future work

- Scalability stress tests
- dApp modularity/composability
- New data acceptance mechanisms
 - Consensus for streamed results

Trustworthy oracles

Practical work

Matthieu BETTINGER

Goal

- Either:
 - Apply trust- and/or event-based logic to your own project
 - Experiment with concepts on a standalone project
 - Trustworthy datafeed mechanisms
 - Request-response infrastructure
 - Staking
 - Consensus:
 - Aggregation (avg, med,...)
 - Voting
 - Random selection
 - ...

Resources

- **Web3J Event listeners**
 - <https://kauri.io/communities/Java%20Ethereum/listening-for-ethereum-smart-contract-events-in-j/>
- **Web3JS Event Listeners**
 - <https://www.pauric.blog/How-to-Query-and-Monitor-Ethereum-Contract-Events-with-Web3/>
- **Solidity events**
 - https://www.tutorialspoint.com/solidity/solidity_events.htm